

# Instruction-Tuning Mistral-7B-Instruct-v0.3 for Automated Essay Scoring and Rationale Generation

Amro Eid, Hossam Shehadeh, Abdullah Jamal

Computer Science Apprenticeship Program, An-Najah National University

## *Summary:*

This project details the instruction-tuning of the Mistral-7B-Instruct-v0.3 model for the sophisticated tasks of Automated Essay Grading (AEG) and interpretable Rationale Generation. Leveraging the Unsloth library for 4-bit quantization (bnb-4bit) and Parameter-Efficient Fine-Tuning (PEFT) techniques, specifically Low-Rank Adaptation (LoRA), the study aimed to create an efficient and effective AEG system. The model was fine-tuned on a custom dataset of 3552 samples, generated via knowledge distillation from a GPT-4o mini teacher model, comprising essay questions, reference answers, student essays, mark schemes, scores, and corresponding rationales. Methodologies included systematic hyperparameter optimization for LoRA configurations (rank, alpha, dropout), iterative dataset balancing, and prompt engineering. The system's performance was evaluated using a comprehensive suite of metrics, including Exact Match Accuracy, Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), Cohen's Kappa, and Weighted F1-score. Findings indicate that higher LoRA capacity ( $r=32$ ,  $\alpha=32$ ) significantly enhances grading accuracy, error reduction, and agreement with human-like grading patterns, demonstrating the potential of instruction-tuned LLMs to provide nuanced, reliable, and explainable feedback in educational assessments.

## I. INTRODUCTION

### **Motivation and Background**

The assessment of student essays is a cornerstone of education, providing valuable insights into comprehension, critical thinking, and writing proficiency. However, the manual grading process is notoriously time-consuming for educators, particularly in large-scale educational settings. Beyond the sheer volume, manual grading can also be prone to subjectivity and inter-rater variability, potentially leading to inconsistencies in evaluation. This burden not only impacts instructor workload but can also delay the provision of timely feedback, which is crucial for student learning and improvement.

To address these challenges, Automated Essay Grading (AEG) systems have been an active area of research for several decades. Early AEG systems often relied on surface-level linguistic features, statistical methods, or rule-based approaches. While these systems demonstrated some success, they frequently struggled with capturing the deeper semantic nuances, coherence, and argumentative structure of complex essays, and often lacked the ability to provide rich, explanatory feedback akin to that of a human grader.

The recent advent of Large Language Models (LLMs) has ushered in a new era for natural language processing, demonstrating remarkable capabilities in understanding, generation, and

reasoning. Specifically, instruction-tuned LLMs, which are fine-tuned to follow natural language prompts and perform specific tasks, offer a promising avenue for developing more sophisticated and nuanced AEG systems. These models can be guided to assess essays based on provided criteria, moving beyond simple keyword matching to a more holistic evaluation.

This project leverages the power of instruction-tuned transformers, specifically the `unsloth/mistral-7b-instruct-v0.3-bnb-4bit` model, to develop an intelligent essay grading system. The Mistral-7B architecture is recognized for its strong performance and efficiency relative to its size. The use of a 4-bit quantized version (`bnb-4bit`) through the Unsloth library is a key aspect, aiming to make fine-tuning and deployment of such powerful models more accessible by significantly reducing computational and memory requirements. This efficiency is further enhanced by employing Parameter-Efficient Fine-Tuning (PEFT) techniques like Low-Rank Adaptation (LoRA), which allows for adaptation of the LLM to the specific task of essay grading by training only a small fraction of the model's parameters.

The goal of this research is to build an AEG system that not only assigns a score but also generates an interpretable rationale for that score, mirroring the detailed feedback an expert human grader might provide. By fine-tuning the model on a custom dataset comprising questions, reference answers, student essays, mark schemes, scores, and rationales, we aim to develop a system that is consistent, scalable, and capable of providing valuable, actionable feedback to students, thereby enhancing the educational assessment process.

## II. DATASET

The foundation of this research is a dataset specifically constructed through a knowledge distillation process, designed to instruction-tune our student model, `unsloth/mistral-7b-instruct-v0.3-bnb-4bit`. This dataset comprises a total of **3552 samples**, where the target outputs (scores and rationales) were generated by a larger, more capable teacher model, **GPT-4o mini**. The aim is to transfer the grading "knowledge" and articulate reasoning capabilities of GPT-4o mini to the more compact Mistral-7B model.

The dataset covers a diverse range of topics, including **Natural Language Processing (NLP), Computer Science, Software Engineering, Big Data, Databases, Data Structures and Algorithms, Information Retrieval, and Machine Learning**. Additionally, it includes examples and questions from **chemistry, football, martial arts, medicine, Palestinian history, programming**, and various **random general knowledge topics**, providing a broad and realistic training corpus to enhance the model's generalization and grading accuracy across multiple domains.

### Dataset Structure and Content:

Each sample in the dataset is a JSON object containing three primary fields: instruction, input, and output. This structure is tailored to the requirements of instruction fine-tuning, where the model learns to generate the output based on the given instruction and contextual input.

1. **instruction:** This field provides a direct, natural language command to the model. For all samples, the instruction is consistent, clearly outlining the desired task:

"Grade the student's answer to the essay question based on the reference answer and the provided mark scheme. Give a score and rationale."

This uniformity helps the model understand the specific task it needs to perform.

2. **input:** This field encapsulates all the necessary contextual information required by the model to perform the grading. It is a multi-line string meticulously structured to include:
  - **Question:** The essay prompt or question that was given to the student.
  - **Reference Answer:** An exemplary or model answer to the question, typically written by an expert. This serves as a benchmark for quality and content.
  - **Student Answer:** The actual essay or response submitted by the student that requires grading.
  - **Mark Scheme:** A set of criteria or a rubric outlining how the student's answer should be evaluated. This typically lists specific points, concepts, or qualities that should be present in the answer and may indicate how marks are allocated.
3. **output:** This field represents the desired generation from the model – the target it learns to produce. It is also a string, formatted to include:
  - **Score:** The numerical score assigned to the student's answer based on the mark scheme and reference answer.
  - **Rationale:** A textual explanation detailing why the specific score was awarded. This rationale should justify the score by referencing the student's answer in relation to the mark scheme.

### III. METHODOLOGY

#### Overview

This research employs a systematic approach to develop an automated essay grading system using instruction-tuned Large Language Models. The methodology consists of six integrated phases designed to achieve expert-level grading performance through iterative improvement and comprehensive evaluation.

## Model Selection and Architecture

### Base Model Configuration

The system utilizes the unsloth/mistral-7b-instruct-v0.3-bnb-4bit model as the foundation architecture. This selection was based on several key criteria:

- **Instruction-following capabilities:** Pre-trained specifically for understanding and responding to natural language instructions
- **Computational efficiency:** 7B parameter size provides strong performance while remaining computationally feasible
- **Memory optimization:** 4-bit quantization through the Unsloth library significantly reduces memory requirements, enabling fine-tuning on consumer-grade hardware
- **Training acceleration:** Unsloth optimization provides enhanced training speed and memory efficiency

### Parameter-Efficient Fine-Tuning (PEFT)

The methodology employs Low-Rank Adaptation (LoRA) as the primary PEFT technique, targeting key transformer components including query, key, value, output, and feed-forward network projections. This approach enables efficient model adaptation while preserving the base model's general knowledge and offers several advantages:

- Computational efficiency through selective parameter updates
- Reduced memory footprint enabling larger effective batch sizes
- Faster training convergence compared to full fine-tuning
- Preservation of base model capabilities while adding task-specific knowledge

## Data Preparation and Training Setup

### Dataset Transformation

The custom dataset was systematically transformed into Alpaca-style instruction-following format to optimize learning efficiency. This transformation process involved:

- Creating consistent instruction templates to guide model understanding
- Consolidating contextual input information (question, reference answer, student answer, mark scheme)
- Formatting target outputs (scores and rationales) for optimal learning
- Ensuring proper tokenization and sequence management for the Mistral-7B architecture

### Dataset Partitioning Strategy

The complete dataset underwent randomized splitting using fixed seeds to ensure reproducibility:

- **Training Set (80%):** Primary model fine-tuning data (2,842 samples)
- **Validation Set (10%):** Hyperparameter tuning and model selection (355 samples)
- **Test Set (10%):** Final unbiased performance evaluation (355 samples)

This partitioning strategy prevents data leakage while ensuring robust model evaluation across all phases.

## Systematic Hyperparameter Optimization

### Multi-Dimensional Parameter Space

A comprehensive hyperparameter search was conducted across multiple dimensions:

#### Learning Rate Optimization:

- Range: [1e-5, 5e-5, 1e-4, 2e-4, 3e-4]
- Systematic evaluation to identify optimal convergence rates

#### LoRA Configuration Variants:

- Rank values: [16, 32, 64, 128]
- Alpha scaling: [16, 32, 64, 128]
- Dropout rates: [0.0, 0.1, 0.2]

#### Training Dynamics:

- Batch sizes: [2, 4, 8] with gradient accumulation
- Learning rate schedulers: [linear, cosine, polynomial]
- Warmup strategies: [5%, 10%, 15% of total steps]

### Model Variant Comparison Framework

Multiple model variants were systematically trained and compared:

- **Conservative Models:** Lower learning rates with higher regularization
- **Aggressive Models:** Higher learning rates with minimal regularization
- **Balanced Models:** Moderate parameters across all dimensions
- **Specialized Models:** Optimized for specific performance aspects

## Iterative Enhancement Process

### Data-Driven Improvement Cycles

The methodology implemented systematic improvement cycles based on performance analysis:

**Cycle 1 - Dataset Balance Optimization:** Analysis revealed score distribution imbalances affecting model performance. Strategic dataset rebalancing through sampling and augmentation improved performance across all score categories.

**Cycle 2 - Advanced Training Strategies:** Implementation of curriculum learning approaches, progressing from clear examples to challenging cases, demonstrated significant improvements in handling edge cases.

**Cycle 3 - Prompt Engineering Optimization:** Multiple instruction formats were systematically tested on subsets of the training data to identify the most effective prompting approach. Different instruction templates were evaluated, including variations with chain-of-thought reasoning, explicit scoring criteria, and structured output formatting. Performance comparison across these formats determined the optimal instruction template, which was then applied to the complete dataset for enhanced accuracy and rationale quality.

## **Comprehensive Evaluation Framework**

### **Multi-Dimensional Assessment**

The evaluation encompasses both quantitative metrics and qualitative analysis:

#### **Quantitative Metrics:**

- Accuracy measures: Exact match, tolerance-based accuracy ( $\pm 1$ ,  $\pm 2$  points)
- Error analysis: Mean Absolute Error, Mean Squared Error, Root Mean Squared Error
- Correlation assessment: Pearson and Spearman correlation with human scores
- Agreement statistics: Cohen's Kappa, weighted F1-scores
- Per-score analysis: Precision, recall, and F1-scores for individual score levels

#### **Qualitative Evaluation:**

- Rationale quality assessment through human evaluation
- Consistency analysis via repeated scoring of identical essays
- Bias detection through systematic scoring pattern analysis
- Edge case performance evaluation on challenging examples

### **Statistical Validation**

Rigorous statistical analysis ensures reported improvements represent genuine advances:

- Paired t-tests comparing model variants
- Effect size calculations (Cohen's d) for improvement magnitude
- Confidence interval estimation for performance metrics
- Multiple comparison corrections for hyperparameter optimization

## **Implementation and User Interface**

### **Training Infrastructure**

The training framework utilized SFTTrainer (Supervised Fine-tuning Trainer) with optimized configurations:

- Sequence length management for long-form text processing
- Memory-efficient training through gradient checkpointing
- Mixed-precision training for computational efficiency
- CUDA-enabled hardware utilization with memory optimization

### **Interactive Demonstration System**

A Gradio-based web interface was developed to demonstrate system capabilities:

- Real-time essay grading with immediate score and rationale generation
- Pre-loaded sample essays for testing
- Custom input support for user-provided scenarios

## Experiment Tracking and Management

### Weights & Biases Integration

Comprehensive experiment tracking was implemented through WandB:

- Complete hyperparameter configuration logging
- Training metrics monitoring (loss curves, learning rates, gradient norms)
- Validation performance tracking on held-out data
- Model artifact storage and version management
- Automated visualization generation for performance comparison
- Systematic experiment organization through tagging and grouping

This methodology provides a systematic, reproducible approach to developing and evaluating automated essay grading systems while ensuring comprehensive performance assessment and continuous improvement capabilities.

## IV. RESULTS

This table summarizes the performance of a language model fine-tuned using LoRA (Low-Rank Adaptation) with different configurations (r=8, 16, 32):

| Metric                             | LoRA r=8, $\alpha$ =8<br>(Base) | LoRA r=16,<br>$\alpha$ =16 | LoRA r=32, $\alpha$ =32<br>(Best) |
|------------------------------------|---------------------------------|----------------------------|-----------------------------------|
| Exact Match Accuracy               | 0.7893                          | 0.7949                     | <b>0.8258</b>                     |
| Tolerance-1 Accuracy ( $\pm 1$ pt) | 0.9972                          | 0.9972                     | 0.9972                            |
| MAE                                | 0.2163                          | 0.2107                     | <b>0.1798</b>                     |
| RMSE                               | 0.4829                          | 0.4770                     | <b>0.4434</b>                     |
| Score Bias (pred - true)           | 0.0084                          | 0.0084                     | <b>0.0225</b>                     |
| Cohen's Kappa                      | 0.7182                          | 0.7259                     | <b>0.7660</b>                     |
| Weighted F1 Score                  | 0.7884                          | 0.7934                     | <b>0.8241</b>                     |
| Pearson Correlation                | 0.9282                          | 0.9300                     | <b>0.9399</b>                     |
| Spearman Correlation               | 0.9221                          | 0.9234                     | <b>0.9333</b>                     |
| Correctly Graded                   | 281 / 356                       | 283 / 356                  | <b>294 / 356</b>                  |
| Over-Graded                        | 40 / 356                        | 39 / 356                   | <b>36 / 356</b>                   |
| Under-Graded                       | 35 / 356                        | 34 / 356                   | <b>26 / 356</b>                   |

### i. Discussion:

The evaluation results clearly demonstrate that increasing the LoRA rank (**r**) and scaling factor (**lora\_alpha**) leads to measurable improvements in the model's grading performance. Specifically, the configuration with **r=32 and lora\_alpha=32** consistently outperforms

the lower-rank models across all major evaluation categories: accuracy, error rates, correlation, and agreement metrics.

### 1. Accuracy Metrics:

- The **exact match accuracy** improves from 0.7893 (r=8) to **0.8258 (r=32)**, indicating that a higher-rank adapter allows the model to better capture the nuances in score prediction, reducing misclassifications. Interestingly, **tolerance-1 accuracy remains stable at 0.9972** across all models, suggesting that while all configurations predict close scores well, only the higher-rank configuration achieves more exact matches.

### 2. Error Metrics (MAE, RMSE):

- There is a **progressive reduction in both MAE and RMSE** with higher LoRA configurations. The MAE drops from 0.2163 to **0.1798**, and RMSE from 0.4829 to **0.4434**, confirming that r=32 produces predictions that are **both more accurate and less variable**.

### 3. Agreement Metrics (Cohen's Kappa, Weighted F1):

- Cohen's Kappa—a robust measure of agreement that accounts for chance—**increases from 0.7182 to 0.7660**, while the **weighted F1-score improves from 0.7884 to 0.8241**. These trends indicate that the model not only performs better overall but also exhibits a stronger agreement with human-like grading behavior as LoRA capacity increases.

### 4. Correlation Metrics (Pearson & Spearman):

- Both **Pearson (linear)** and **Spearman (rank)** correlation coefficients improve with higher-rank settings. The model at r=32 achieves **0.9399 Pearson and 0.9333 Spearman**, highlighting its capability to consistently align with the actual score distributions and rank ordering of responses.

### 5. Bias and Grading Errors:

- Score bias remains minimal across all settings, but **r=32 shows the fewest over-graded and under-graded samples**, reinforcing that larger LoRA configurations reduce systematic grading inconsistencies.



## ii. Key Observations:

- **Higher LoRA capacity ( $r=32$ ,  $\alpha=32$ ) significantly improves grading quality** across all metrics—accuracy, error, agreement, and correlation.
- **Exact match accuracy improves by  $\sim 4.6\%$**  when comparing  $r=8$  to  $r=32$ , a notable gain in educational contexts where precision in scoring is critical.
- **MAE and RMSE reductions reflect improved prediction stability**, with  $r=32$  yielding the lowest average prediction error.
- **Cohen’s Kappa  $> 0.76$**  for  $r=32$  signifies **substantial agreement** with human annotations, making it suitable for high-stakes automated assessment tasks.
- **Weighted F1 score increases to 0.8241**, reflecting **balanced performance across all score classes**, even with skewed class distributions.
- **Grading fairness improves**:  $r=32$  leads to **fewer over-graded ( $\downarrow 4$  samples) and under-graded ( $\downarrow 9$  samples)** responses compared to the base model.

## V. CONCLUSION

This project successfully demonstrates the efficacy of instruction-tuning the **Mistral-7B-Instruct-v0.3** model, enhanced with **4-bit** quantization via Unsloth and Low-Rank Adaptation (**LoRA**), for automated essay scoring and rationale generation. The systematic fine-tuning process on a custom, knowledge-distilled dataset allowed the model to achieve strong performance in both assigning accurate scores and providing coherent, contextually relevant justifications for those scores.

The results clearly indicate that increasing LoRA capacity (specifically,  $r=32$ ,  $\alpha=32$ ) leads to significant improvements across key evaluation metrics. The model achieved an Exact Match Accuracy of **0.8258**, a low Mean Absolute Error of **0.1798**, a Cohen's Kappa of **0.7660** signifying substantial agreement, and a Weighted F1 Score of **0.8241**. These figures underscore the model's capability not only for precise score prediction but also for generating outputs that align closely with nuanced human assessment. The reduction in grading errors (over-graded and under-graded samples) with optimized LoRA configurations further points to increased grading fairness and consistency.

The developed system offers a promising advancement in Automated Essay Grading, moving beyond simple score assignment to provide valuable, interpretable feedback akin to human graders. This work contributes to the development of more sophisticated, efficient, and equitable educational assessment tools, highlighting the potential of PEFT techniques to adapt powerful LLMs for specialized tasks with reduced computational overhead. Future work could explore further enhancements in rationale quality, broader domain applicability, and integration into real-world educational platforms.

#### REFERENCES

- [1] <https://huggingface.co/mistralai/Mistral-7B-Instruct-v0.2>
- [2] <https://huggingface.co/unsloth/mistral-7b-instruct-v0.2-bnb-4bit>
- [3] <https://docs.unsloth.ai/get-started/fine-tuning-guide>
- [4] ChatGPT , Claude and Google AI Studio

#### AUTHORS

**First Author** – Amro Eid,

B.Sc. Student in Computer Science (Senior Year), Department of Computer Science Apprenticeship Program, NNU, Palestine  
and [amro.eidd@gmail.com](mailto:amro.eidd@gmail.com)

**Second Author** – Hossam Shehadeh,

B.Sc. Student in Computer Science (Senior Year), Department of Computer Science Apprenticeship Program, NNU, Palestine  
and [hossam.h.shehadeh@gmail.com](mailto:hossam.h.shehadeh@gmail.com)

**Third Author** – Abdullah Jamal,

B.Sc. Student in Computer Science (Senior Year), Department of Computer Science Apprenticeship Program, NNU, Palestine  
and [abood.jamal005@gmail.com](mailto:abood.jamal005@gmail.com)